
Heap Sort

Heaps, Heap property, Building a heap, Heap sort algorithm

Rashed Hassan Siam

Department of CSE

Eastern University

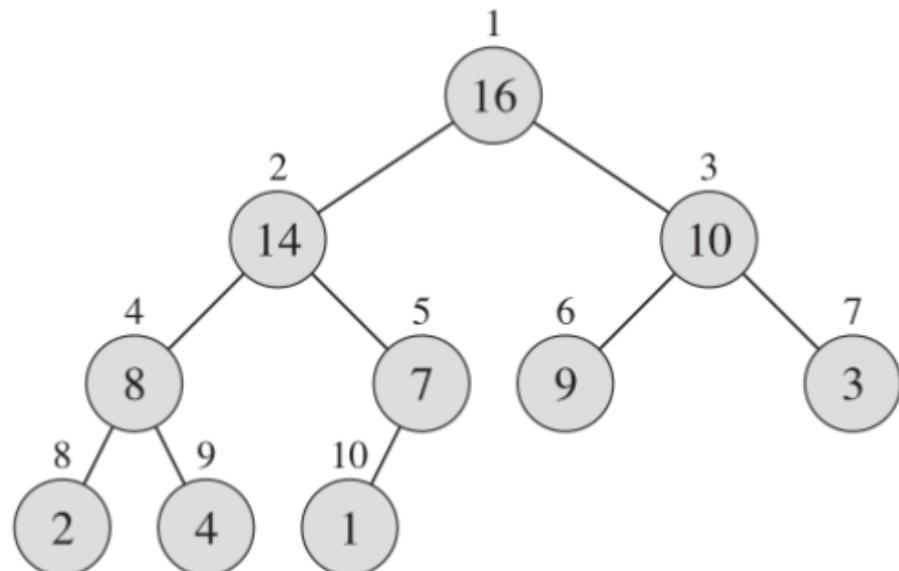
Introduction to Heaps

- Heapsort is a sorting algorithm with a running time of $O(n \log n)$.
- Like insertion sort, it sorts **in place**. No auxiliary memory space is needed. So, space complexity is $O(1)$.
- It achieves this using a specific data structure called a "heap".

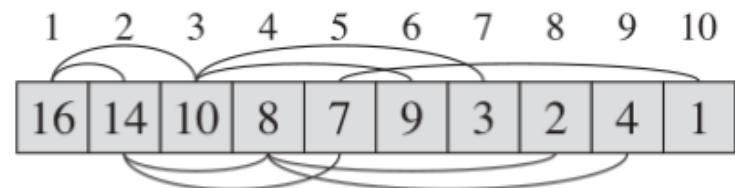
The Heap Data Structure

- The (binary) heap data structure is an array object that can be viewed as a nearly complete binary tree.
 - Each node in the tree corresponds to an element in the array.
 - The tree is completely filled on all levels, except possibly the lowest level, which is filled from the left.
 - The array object, A, has two attributes: A.length (the number of elements in the array) and A.heap-size (how many elements are currently valid elements of the heap).
 - The root of the tree is A[1].

The Heap Data Structure



(a)



(b)

A max-heap viewed as (a) a binary tree and (b) an array.

The Heap Data Structure

- Given the index i of a node, we can compute the indices of its parent and children:

PARENT(i)

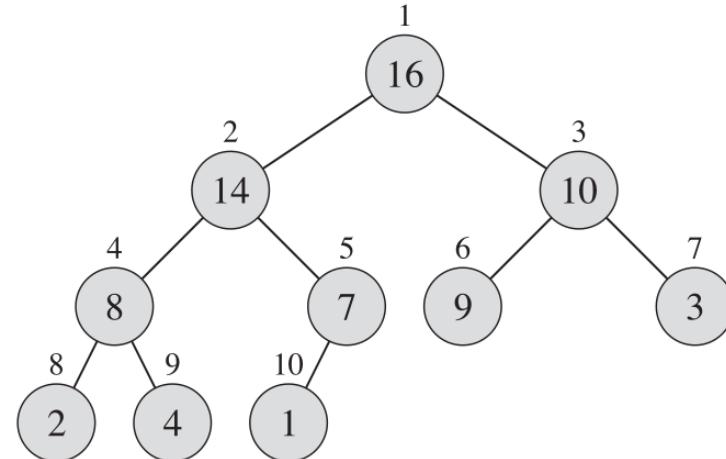
```
1 return  $\lfloor i/2 \rfloor$ 
```

LEFT(i)

```
1 return  $2i$ 
```

RIGHT(i)

```
1 return  $2i + 1$ 
```



- A heap of n elements has a height of $\Theta(\log n)$.

The Heap Property

- There are two kinds of binary heaps: max-heaps and min-heaps. In both cases, the values in the nodes must satisfy a specific "heap property".
 - Max-Heaps
 - In a max-heap, the value of a node is at most the value of its parent. This is the max-heap property: $A[\text{PARENT}(i)] \geq A[i]$ for every node i other than the root.
 - This means the largest element in a max-heap is stored at the root ($A[1]$). The heapsort algorithm uses max-heaps.

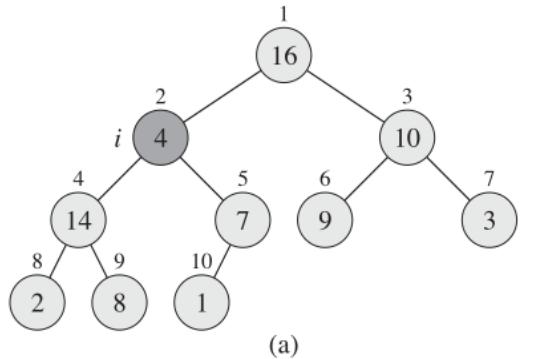
The Heap Property

- There are two kinds of binary heaps: max-heaps and min-heaps. In both cases, the values in the nodes must satisfy a specific "heap property".
 - Min-Heaps
 - A min-heap is organized in the opposite way. The min-heap property is: $A[\text{PARENT}(i)] \leq A[i]$ for every node i other than the root.
 - In a min-heap, the smallest element is at the root. Min-heaps are commonly used to implement priority queues.

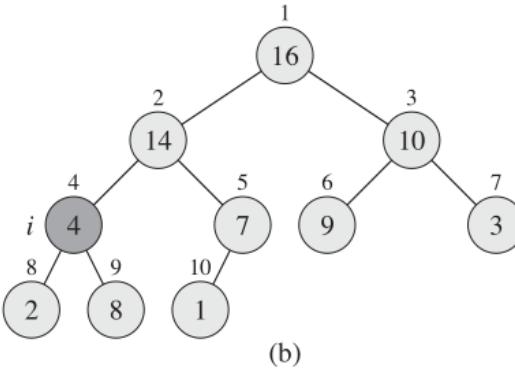
Maintaining the Property: MAX-HEAPIFY

- To maintain the max-heap property, we use the MAX-HEAPIFY procedure.
 - Assumption: When $\text{MAX-HEAPIFY}(A, i)$ is called, it assumes the binary trees rooted at $\text{LEFT}(i)$ and $\text{RIGHT}(i)$ are already max-heaps.
 - Problem: $A[i]$ might be smaller than its children, violating the property.
 - Action: MAX-HEAPIFY lets the value at $A[i]$ "float down" in the max-heap so that the subtree rooted at index i obeys the max-heap property.

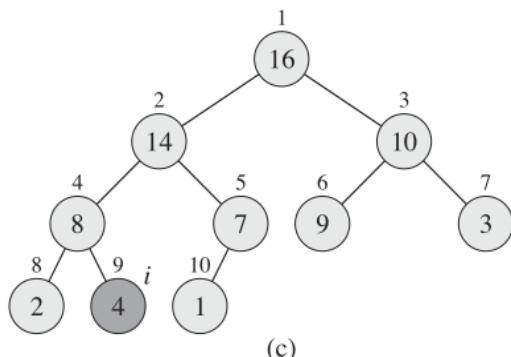
Maintaining the Property: MAX-HEAPIFY



(a)



(b)



(c)

MAX-HEAPIFY(A, i)

```
1   $l = \text{LEFT}(i)$ 
2   $r = \text{RIGHT}(i)$ 
3  if  $l \leq A.\text{heap-size}$  and  $A[l] > A[i]$ 
4       $\text{largest} = l$ 
5  else  $\text{largest} = i$ 
6  if  $r \leq A.\text{heap-size}$  and  $A[r] > A[\text{largest}]$ 
7       $\text{largest} = r$ 
8  if  $\text{largest} \neq i$ 
9      exchange  $A[i]$  with  $A[\text{largest}]$ 
10     MAX-HEAPIFY( $A, \text{largest}$ )
```

The action of MAX-HEAPIFY($A, 2$), where $A.\text{heap-size} = 10$.

Maintaining the Property: MAX-HEAPIFY

- This procedure works by finding the largest of the elements $A[i]$, $A[LEFT(i)]$, and $A[RIGHT(i)]$.
- If $A[i]$ is not the largest, it is swapped with the largest child. This swap may cause the subtree rooted at the child (now at index `largest`) to violate the max-heap property, so `MAX-HEAPIFY` is called recursively on that subtree.
- The running time of `MAX-HEAPIFY` on a node of height h is $O(h)$, or $O(\log n)$ on a heap of size n .

Building a Heap

- We can build a max-heap from an unordered input array $A[1..n]$ by using the MAX-HEAPIFY procedure in a bottom-up manner.
- The elements in the subarray $A[(\lfloor n/2 \rfloor + 1) .. n]$ are all leaves of the tree, so they are already trivial 1-element heaps.
- The BUILD-MAX-HEAP procedure goes through the remaining (non-leaf) nodes and runs MAX-HEAPIFY on each one.

Building a Heap

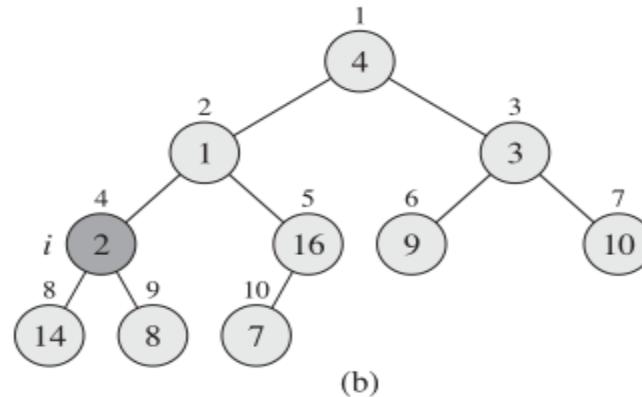
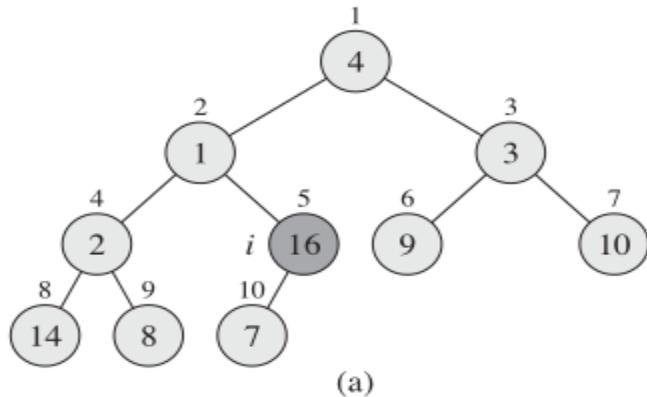
- This works because of a key loop invariant:
 - At the start of each iteration of the for loop (for index i), each node $i+1, i+2, \dots, n$ is the root of a max-heap.
- When MAX-HEAPIFY(A, i) is called, the loop invariant ensures that the children of node i (which are numbered higher than i) are already roots of max-heaps, which is the condition required for MAX-HEAPIFY to work correctly.
- At termination, $i=0$, and node 1 is the root of a max-heap for the entire array.

BUILD-MAX-HEAP(A)

```
1  $A.\text{heap-size} = A.\text{length}$ 
2 for  $i = \lfloor A.\text{length}/2 \rfloor$  downto 1
3   MAX-HEAPIFY( $A, i$ )
```

Building a Heap

A	4	1	3	2	16	9	10	14	8	7
-----	---	---	---	---	----	---	----	----	---	---

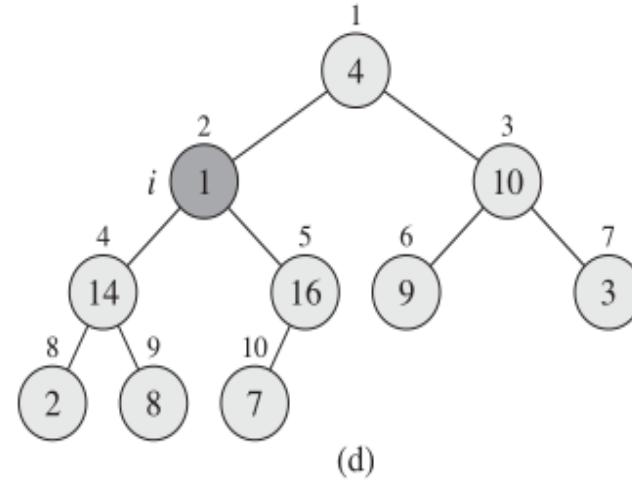
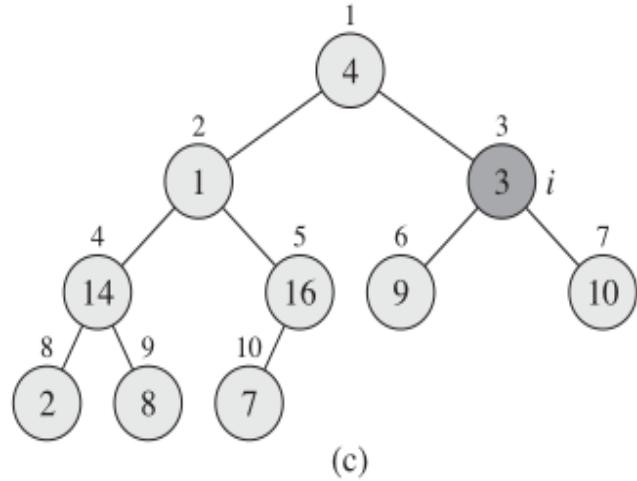


The operation of BUILD-MAX-HEAP, showing the data structure before the call to MAX-HEAPIFY in line 3 of BUILD-MAX-HEAP.

(a) A 10-element input array A and the binary tree it represents. The figure shows that the loop index i refers to node 5 before the call $\text{MAX-HEAPIFY}(A, i)$.

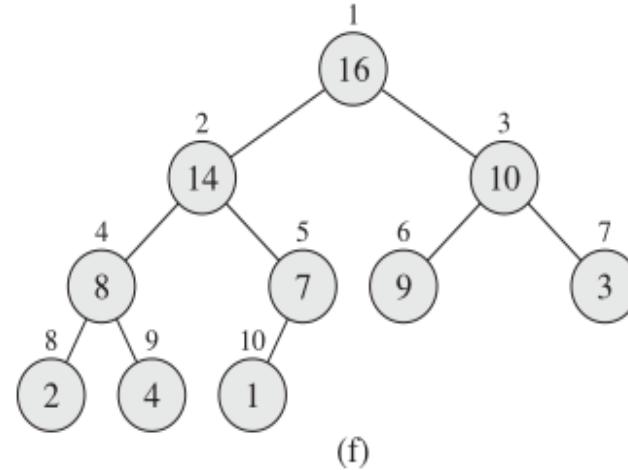
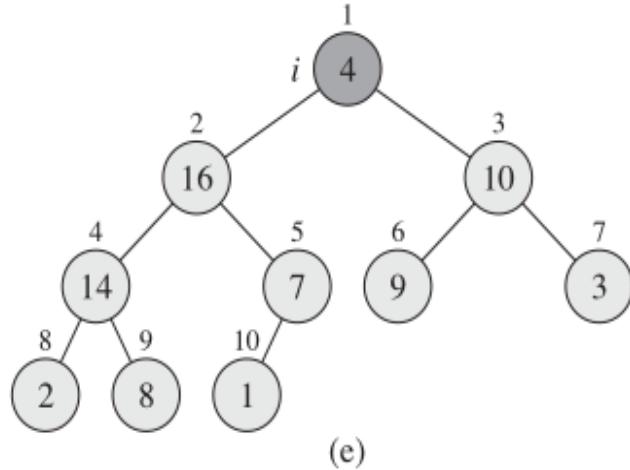
(b) The data structure that results.

Building a Heap



(c)-(d) Subsequent iterations of the for loop in BUILD-MAX-HEAP.

Building a Heap



- (e) Subsequent iteration of the for loop in BUILD-MAX-HEAP.
(f) The max-heap after BUILD-MAX-HEAP finishes.

Building a Heap

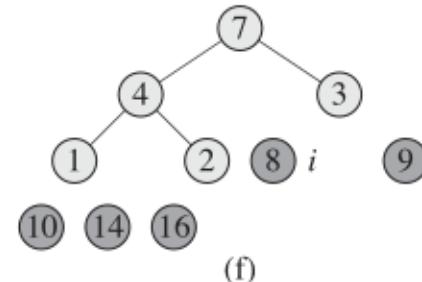
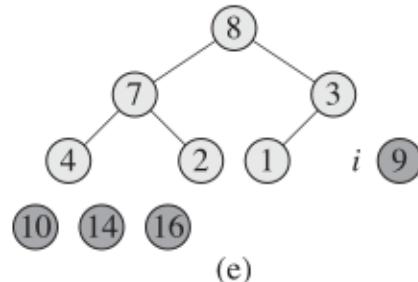
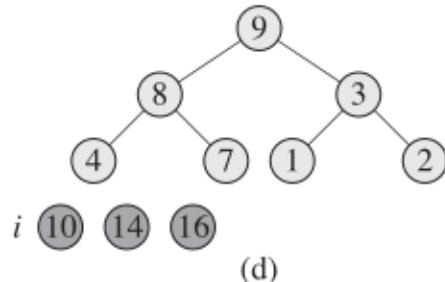
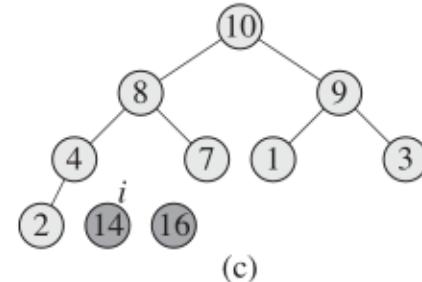
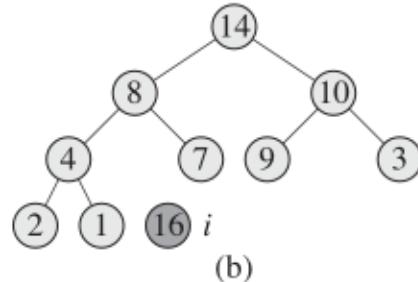
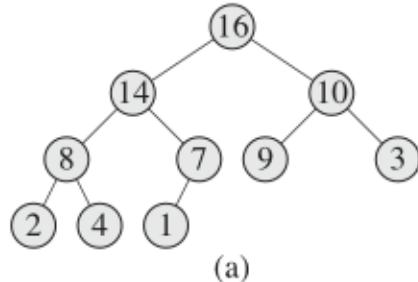
- We can compute a simple upper bound on the running time of BUILD-MAX-HEAP as follows:
 - Each call to MAX-HEAPIFY costs $O(\log n)$ time, and BUILD-MAX-HEAP makes $O(n)$ such calls.
 - Thus, the running time is $O(n \log n)$. However, this bound is not asymptotically tight.
- A tighter analysis shows that the time required by MAX-HEAPIFY varies with the height of the node, and most nodes have small heights (A max-heap of 1073741823 elements has a very small height of 29 only!).
 - So, the number of levels $\log n$ can be omitted as a lower order term.
 - Thus, we can bound the running time of BUILD-MAX-HEAP as $O(n)$.
 - Hence, we can build a max-heap from an unordered array in linear time.

The Heapsort Algorithm

The heapsort algorithm combines BUILD-MAX-HEAP and MAX-HEAPIFY to sort an array.

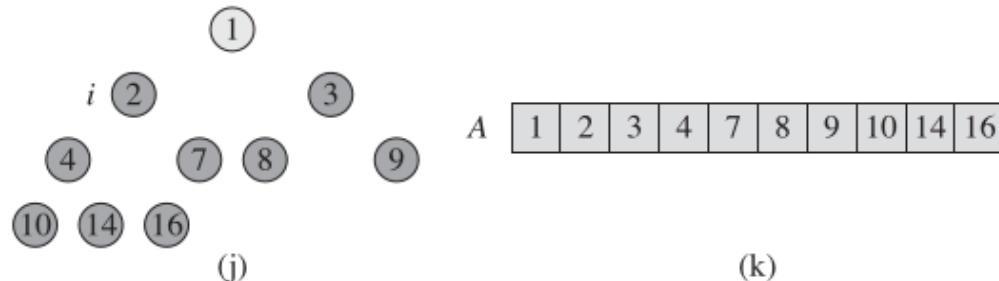
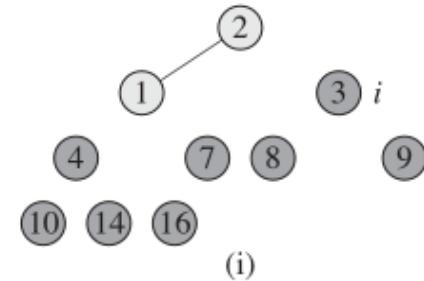
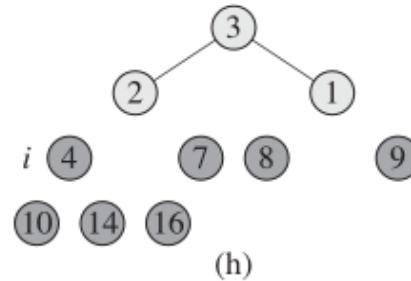
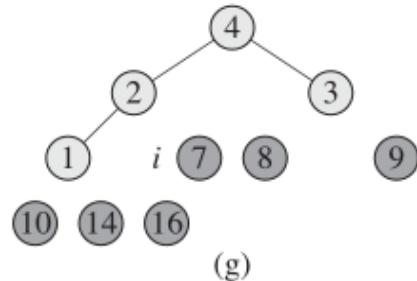
1. It starts by using BUILD-MAX-HEAP to build a max-heap on the input array $A[1..n]$.
2. Since the maximum element of the array is now at the root $A[1]$, it is exchanged with $A[n]$, placing the largest element in its correct final position.
3. The heap size is reduced by one, "discarding" node n from the heap.
4. The new root element might violate the max-heap property. To restore it, $\text{MAX-HEAPIFY}(A, 1)$ is called, which leaves a max-heap in $A[1..n-1]$.
5. This process (steps 2-4) is repeated for the max-heap of size $n-1$ down to a heap of size 2.

The Heapsort Algorithm



The operation of HEAPSORT

The Heapsort Algorithm



The operation of HEAPSORT

The Heapsort Algorithm

- Running Time: The HEAPSORT procedure takes $O(n \log n)$ time.
 - The initial call to BUILD-MAX-HEAP takes $O(n)$ time.
 - There are $n-1$ calls to MAX-HEAPIFY (one for each iteration of the loop).
 - Each call to MAX-HEAPIFY takes $O(\log n)$ time.
 - Thus, the total time is $O(n) + O((n-1) \log n)$, which is $O(n \log n)$.

HEAPSORT(A)

```
1 BUILD-MAX-HEAP( $A$ )
2 for  $i = A.length$  down to 2
3   exchange  $A[1]$  with  $A[i]$ 
4    $A.heap-size = A.heap-size - 1$ 
5   MAX-HEAPIFY( $A, 1$ )
```

Reference

- *Introduction to Algorithms, Third Edition*
 - *Thomas H. Cormen, Charles E. Leiserson, Ronald L. Rivest, Clifford Stein*